ALEX CHEN

alex.chen@example.com · (555) 010-0142 · San Francisco, CA · github.com/example · example.com

SUMMARY

AI engineer with a statistics and distributed-systems background, 3 years taking LLM features from prototype to production: retrieval, agents, and the evals that keep them honest. M.S. @ Lakeside University.

EDUCATION

Lakeside University, M.S. in Computer Science — Aug 2022 – May 2024

GPA 3.9/4.0 · Distributed Systems · Information Retrieval · Machine Learning · Teaching Assistant for Databases

Riverside Institute of Technology, B.S. in Statistics — Sep 2018 – Jun 2022

GPA 3.8/4.0 · Dean's List (6 semesters) · Undergraduate research in Bayesian time-series forecasting

WORK EXPERIENCE

Northwind Labs — Jul 2024 – Present

AI Engineer — San Francisco, CA

- Built a tool-calling agent runtime with retries and idempotent steps, raising end-to-end task success from 61% to 84%
- Designed an offline eval suite of 1,200 graded traces with an LLM judge, catching 9 regressions before they shipped
- Cut median agent latency 38% by parallelizing independent tool calls and caching retrieval results per conversation
- Moved 14 internal tools behind a shared MCP server, so product teams add agent capabilities without runtime changes
- Introduced a versioned prompt registry with staged rollouts and A/B gates, halving the time to ship a prompt change

Helio Health — May 2023 – Aug 2023

Machine Learning Intern — Remote

- Shipped hybrid BM25 and dense retrieval over 2M clinical notes, lifting top-5 recall from 71% to 89% on a labeled set
- Fine-tuned a cross-encoder reranker with LoRA on 40k physician queries, improving answer precision by 14 points
- Added a PHI redaction pass using spaCy rules and a small NER model, clearing HIPAA review before the hospital pilot

Cobalt Analytics — Jun 2021 – Jul 2022

Software Engineer — Austin, TX

- Migrated nightly batch ETL to Kafka streaming jobs, cutting data freshness from 24 hours to under 10 minutes for clients
- Added data-quality checks and on-call alerting in Grafana, reducing customer-reported data incidents by 70% in 2 months
- Rewrote the slowest dashboard queries as PostgreSQL materialized views, dropping p95 page load time from 9s to 1.2s

Riverside Institute of Technology — Jan 2020 – May 2021

Undergraduate Research Assistant — Riverside, CA

- Built a hierarchical Bayesian model in Stan for regional power demand, beating the ARIMA baseline by 18% on MAPE
- Co-authored a workshop paper on calibrated uncertainty for short-term load forecasts, presented at a 2021 symposium
- Packaged the forecasting pipeline as an open-source R library with 40 unit tests, reused by two later student theses

PROJECTS

TraceLens ↗ — Open-source viewer for agent traces with diffing across runs and per-step cost breakdowns

Quill ↗ — Writing assistant that grounds every suggestion in cited sources, built on a local vector store

PromptDiff ↗ — CLI that regression-tests prompts against golden outputs in CI, used by three open-source projects

TALKS & AWARDS

Speaker, Bay Area LLMOps Meetup (2025): "Evals before vibes: grading agent traces with an LLM judge"

First place, Lakeside University AI Hackathon (2023) · Best Use of Retrieval, Open Data Hack Weekend (2022)

Volunteer mentor, Code Forward fellowship (2024 – present): coaching first-generation students through ML projects

SKILLS

Languages	Python, TypeScript, Go, SQL, Bash
Infrastructure	AWS, Docker, Kubernetes, Terraform, Kafka, PostgreSQL, Redis
Data	pandas, NumPy, scikit-learn, Stan, Spark, Airflow, dbt
AI & LLMOps	PyTorch, LangGraph, vLLM, LoRA, RAG, LLM-as-judge evals, OpenTelemetry